

Current measured research architectures · future scale-up remains unstated

DenseCompute

9 layers · width 320 · context 1024
16.57M stored / 16.57M active (estimate)
Dense feed-forward blocks

SparseV3

9 layers · width 320 · context 1024
21.56M stored / 16.49M active (estimate)
Grouped Top-1 · 12 experts at layers 2, 5, 8

SparseV3 + Ngram

9 layers · width 320 · context 1024
25.77M stored / 16.50M active (estimate)
Grouped Top-1 · 12 experts at layers 2, 5, 8
4 hashed Ngram tables · layer 1 residual

Shared tokenizer, real corpus, token order, seed42, batch and 100M LR schedule